Exercise 3: LLVM Analysis & Transformation

In this assignment, you should implement two LLVM analysis passes that detect, print and correct all variable uses before initialization. The assignment is divided in two parts: first, you should be able to detect and print (on the LLVM error stream named `errs`) all uses before initialization. Subsequently, you should be able to fix the input IR by adding an initialization instruction for each uninitialized variable.

You must develop these analysis tools as two LLVM compiler passes implemented in a single `*.cpp` file. A pass template file (`pass.cpp`) is given in order to implement your solution.

LLVM and C++ Basics

LLVM is written in C++. Even if you are not familiar with it, to implement an LLVM pass only few language features are necessary. Should you have any doubts, browse the examples included in the LLVM examples archive, check the online documentation http://llvm.org/docs and in particular the LLVM programming guide http://llvm.org/docs/ProgrammersManual.html. The programming guide briefly describes important data-structures and shows how common tasks (like inserting an instruction) can be accomplished.

An LLVM pass is a derived class (subclass) of the `Pass` class, which implements the analysis by overriding virtual methods inherited from `Pass`. The archive code has many examples, including BasicBlock, Function, Loop, Module, Region as well as SCC passes. For example, you can implement a Function pass iterating all functions in the input code by implementing a class that derives (extends) the base class `FunctionPass` and overriding the virtual function `bool runOnFunction(Function &F)`. A pass implementation should be registered to the pass manager with a `RegisterPass<MyPass>` template function (see http://llvm.org/docs/WritingAnLLVMPass.html).

LLVM uses C++ reference (e.g., `Function &F`), which are used like normal values but conceptually behave as a pointer (https://www.dgp.toronto.edu/~patrick/csc418/wi2004/notes/PointersVsRef.pdf).

C++ provides a Standard Template Library (STL) offering a number of useful containers and algorithms, such as list, vector and map (see https://www.sgi.com/tech/stl/stl_introduction.html). LLVM additionally provides a number of more specialized containers such as `SmallSet` and `SetVector`. Maps are particularly useful to store elements associated with a key element (e.g., you can associate a struct with a basic block) http://www.cplusplus.com/reference/map/map ; LLVM provides specific implementations for `StringMap`, `IndexedMap`, `DenseMap`, `ValueMap`, and `IntervalMap`.

LLVM instructions all extend the base class `Instruction`, which in turn extends `Value`. For this assignment the `AllocaInst`, `LoadInst` and `StoreInst` instructions are particularly important. The LLVM `dyn_cast<>` or C++'s `dynamic_cast<>` can be used to check whether a pointer/reference is a valid subclass pointer/reference of another class. For example, a dynamic cast can be used to check whether an `Instruction` is an `AllocaInst`.

To print a string you should use the predefined LLVM raw error stream:

```
errs() << "This is a message on LLVM error stream" << end;
```

Part 1: Analysis Pass [Read only]

You must develop an LLVM pass that detects and prints possible uses before initialization, therefore it should not change the input IR.

Your code must run on the IR generated by Clang in debug mode (-g -O0) and **should not** apply any other LLVM passes. You are not allowed to use the PassManager (expect for the pass registration, see below). As a consequence, the IR you will work on **is not** in SSA form (i.e., no PHI nodes).

As in the provided template file, you should register your pass with the following code:

```
static RegisterPass<DefinitionPass> X("def-pass", "Uninitialized
variable pass");
```

Here `DefinitionPass` is the name of the class implementing the analysis pass, and "def-pass" is the command line argument we will use later to activate the pass from the optimizer (`opt`).

While there are multiple ways to implement this analysis, we strongly recommend an implementation using iterative data-flow analysis. Some suggestions on how to approach the problem:

1. Solve the problem on paper. Determine what the direction of the analysis is, what the gen and kill sets are, how the data-flow equations for in/out look like, which meet to use (union or intersection) and how the sets need to be initialized. Because the analysis is very simple, you may find that either the gen or the kill set is always empty.

2. It is customary to represent the gen/kill/in/out sets as bitvectors. In LLVM the `BitVector` class can be used for this purpose. In order to use a bitvector, you need to establish a mapping between variables (or LLVM `Value`) and bits. As such, the first step is to create a map (e.g. `ValueMap`) from the relevant values to bit offsets. Similarly, you will also need to create a mapping between `BasicBlock`s and a structure containing the gen/kill/in/out sets for each block.

3. Next, compute the gen and kill sets for each block. Then, implement the equations for in/out. The `BitVector` class provides the necessary set operations (union, intersection, complement). To access the successors or predecessors of a basic block (depending on analysis direction) a `pred_iterator` or `succ_iterator` can be used. You can either continue recomputing all blocks until a fixed point is reached, or use a worklist to only recompute blocks that might have changed.

4. Finally, print all uses of potentially uninitialized variables using `getName()` and `getDebugLoc()`.

Part 2: Fixing Pass [Modify IR]

The second pass will fix the code by adding an initialization instruction for each non-initialized variable. The initialization value is

- 10 for integer
- 20.0 for float
- 30.0 for double

to be used whenever a variable is used before initialization. Therefore, the second pass will actually change the IR.

To check the type of a value you can use `Type::isIntegerTy()`, `Type::isFloatTy()` or `Type::isDoubleTy()`. To create a new store instruction in the IR, you should create a `StoreInst` object (e.g., `new StoreInst(…)`); then, you can use either a constructor parameter or the method `insertAfter()` to add it in the right place.

Limitations

For this assignment, you can only use the IR produced by Clang, as-is. You cannot use other LLVM passes or invoke the LLVM pass manager to call other external code.

As the required pass is machine independent, it will be invoked by `opt` (as we have seen in the other examples). You are not allowed to call LLVM's internal analyses such as dominator tree or live interval construction, but you can implement your own within the submitted code, if needed.

In some cases, a constant propagation pass may further improve the analysis by skipping some paths in the CFG, but this **is not** required in this assignment (and the given expected output does not implement it).

For this assignment you only need to consider instructions of type `AllocaInst`, `LoadInst` and `StoreInst`. It is not necessary to take more complicated situations (such as taking the address of a local variable using `&var` and indirectly initializing it) into account.

We use LLVM 3.5 for evaluation. If you use a different LLVM version for development, make sure your code works with LLVM 3.5 prior to submission.

Evaluation

For the first pass, the **error stream** of your LLVM def pass (for each input code) should closely match the one provided by our implementation (`test*.def`), which prints at which points in the program a variable may be used prior to initialization (an uninitialized variable can also be used twice in one line, in which case you should also print two messages).

For the second pass, instead, the **output stream** of the *transformed* code (for each input code, after applying the fix pass) executed with the LLVM JIT (i.e., the command `lli`) will be compared with our implementation (`test*.fix`).

You can use the provided `./test.sh` script to check whether your results are correct.

Your C++ file should also include a brief explanation of the algorithm you used to solve this assignment at the top.

Getting started

To start, you have a file named `pass.cpp` with an empty implementation of the def and fix passes. Remember to rename `pass.cpp` with your file name, and to apply this change also in the given `Makefile`. The `./test.sh` script checks your pass against the expected result. The following files are produced:

- testN.c: The tested C file.

- testN.bc: LLVM bitcode file as generated by Clang.
- testN.ll: Disassembled LLVM IR of the bitcode file.
- testN.def: The expected output from the definition pass (potentially uninitialized variables).
- testN.def.out: The output your definition pass produced.
- testN.def.diff: Difference between your output and the expected output.
- testN.fix: The expected output of the fixed code (after running the fixing pass).
- testN.fix.bc: The bitcode produced by your fixing pass.
- testN.fix.out: The output your fixed code produces.
- testN.fix.diff: Difference between your output and the expected output.

You can check the `./test.sh` script to see how the files can be manually generated, and how the necessary utilities (clang, opt, lli, llvm-dis) are invoked.

Submission
- Use the ISIS website
- Submit the project as a single file, extending the given sample file
- File name has to be: `lastname1_lastname2.cpp`, for example: `maradona_klinsmann.cpp`
- First line of the `cpp` file should include first name, surname and student id, for each group participant, e.g.
  `/* Diego Maradona 10, Juergen Klinsmann 18 */`
- A brief explanation of the algorithm you used to solve this assignment should also be included at the top of the file.
- The two passes must be registered, respectively, as `def-pass` and `fix-pass`, as shown in the input file